

A Frequency Dial for Structure-Conditioned Diffusion

Anonymous submission

Abstract

A structure-conditioned diffusion adapter is trained for one operating point. Canny edges, a depth map, a sketch: each fixes how tightly the output must follow the condition, and moving that trade-off means training again. We train a single adapter on conditions whose information content is randomised — a high-pass filter whose cutoff r is drawn per sample — and read r back at test time as a dial. On FFHQ-512, structure consistency runs from 0.7469 at $r=0.05$ to 0.2347 at $r=0.3$ while sample diversity rises from 0.3126 to 0.3794, and FID stays between 51.58 and 54.68 over that span. We report the sweep out to $r=1$ but claim the dial only on $r \in [0, 0.3]$, above which the condition retains too little energy for structure consistency to be interpretable. Turning the adapter’s strength down is not the same knob. At matched structure consistency — scale 0.50 reaches 0.2146 and the cutoff $r=0.3$ reaches 0.2347 — the strength setting costs FID 76.46 where the cutoff costs 52.77, so the cutoff dominates on both axes at once. Once an artist filters a hand drawing at r , the dial works on drawings too.

Two findings came out of building it. The frozen SD VAE destroys a high-pass condition before any adapter can read it: only 0.5773 of the condition survives an encode–decode round trip at $r=0.3$ and 0.0700 at $r=0.7$. And between these two injection designs placement matters more than width — skip-only retains SC 0.5370 against 0.2112 for attention-only, with both paths at 0.6991 — while quadrupling the adapter’s width inside the attention-only design buys far less. The adapter trains 6,611,469 parameters, small next to ControlNet but not next to recent lightweight control modules; we make no efficiency claim.

1 Introduction

Image editing has long given users a scale or frequency band to turn [5, 28], and sketch-based face synthesis has treated a drawing as a soft constraint rather than a template [1], in some systems with an explicit control over how far the sketch is trusted [46]. ControlNet [50] and its successors [19, 25, 34, 53] moved that control into training. An adapter trained on Canny edges follows Canny edges; one trained on depth follows depth. Following a structure more loosely then means picking a different checkpoint or scaling the adapter down at inference.

We change what the condition is rather than how hard the model listens to it. Multi-scale editing removes or amplifies a band of the output image [5, 28]; here the band is removed from the condition, and the generative prior decides what fills it. The condition is the image after a radial high-pass filter with normalised cutoff r . At $r=0$ the model receives the grayscale image and the task is colourisation.

As r grows the filter removes more of the spectrum, until at $r=1$ almost nothing is left: measured on 500 FFHQ test images, a cutoff of 0.05 already discards 97.5 % of the spectral energy. Training draws r per sample, so one set of weights sees the whole range, and r becomes an inference-time parameter.

Our contributions are:

- A control dial. One adapter spans structure consistency 0.7469 to 0.2347 on $r \in [0.05, 0.3]$ with FID between 51.58 and 54.68 (Sec. 4.2).
- Evidence that the dial is not a rescaled version of adapter strength. At the closest matched pair the two settings reach structure consistency 0.2146 and 0.2347, at FID 76.46 and 52.77 (Sec. 4.3).
- A measurement of how much of a high-pass condition survives the frozen VAE, which explains why latent-space conditioning fails for sparse structure (Sec. 4.4).
- A controlled comparison of injection sites showing that placement matters more than adapter capacity or receptive field (Sec. 4.5).

2 Related work

Conditional control. Text-to-image diffusion [13, 26, 36, 37, 39] became spatially controllable with ControlNet [50], which reads the condition image with a small convolutional network, duplicates the encoder and injects residuals through zero convolutions. T2I-Adapter [25] does the same job with a small feature extractor, IP-Adapter [47] adds decoupled cross-attention for image prompts, and GLIGEN [20] grounds generation on boxes. Uni-ControlNet [53] and UniControl [34] fold many condition types into one model, and ControlNet++ [19] closes a consistency loop on the produced image. A second line shrinks the control branch on the same UNet: ControlNetXS [49] couples a much smaller control network to encoder and decoder in both directions, CtrLoRA [45] adapts a shared base ControlNet to each new condition with LoRA, and Simple-ControlNet [6] feeds the condition through a single-scale projection at the top of the UNet with LoRA on self-attention, in under 3M parameters, exposing a scale on that projection at inference. Each of these fixes the condition type at training and, with it, how tightly the output follows; an inference-time scale such as Simple-ControlNet’s moves that operating point by weakening the adapter. Our cutoff moves it through the condition on one set of weights, and on our adapter the two controls separate in FID (Sec. 4.3).

Efficient control on transformer backbones. As backbones moved from UNets to diffusion transformers [4, 30],

control moved from duplicated encoders to lightweight modules. ControlNeXt [31] removes the parallel branch, OminiControl [40] reuses the backbone itself as the condition processor, and EasyControl [52] and NanoControl [22] inject conditions through LoRA-style [14] modules at a small fraction of the parameter cost. Recent work also revisits how conditions enter these models: Laytrol [16] inherits the backbone’s parameters to preserve pretrained knowledge under layout control, PixelControl [21] notes that the VAE’s spatial compression can weaken high-frequency and low-area condition signals, and Zhou et al. [54] replace concatenate-and-attend with position-aligned and keyword-scoped attention in multi-condition DiTs. On a UNet the question has a concrete form, since skip tensors are available as an entry point: in our comparison (Sec. 4.5) skip residuals carry more structure than cross-attention.

Frequency bands as an editing axis. Band decompositions are a standard editing axis in graphics: detail manipulation rescales the bands of an edge-preserving or Laplacian decomposition [5, 28], and hybrid images combine the low spatial frequencies of one picture with the high frequencies of another, so that the interpretation changes with viewing distance [27]. In diffusion, frequency decomposition appears mostly as an intervention on the model’s own signals rather than on its condition. FreeU [38] rebalances skip and backbone features spectrally at no training cost, FreSca [15] scales frequency bands of the noise prediction using spatial or energy-based cutoffs, DeCo [23] decouples frequencies in pixel-space diffusion, and FeRA [48] routes adaptation by frequency energy. FOCUS [41] conditions the reverse process on learned frequency priors for test-time adaptation. Closest to us in the choice of control, frequency bands also act as a knob on image-to-image guidance: FCDiffusion [8] extracts DCT bands of the source features and switches among frequency control branches, FBSDiff [7] sets the guiding intensity of a reference through the bandwidth of a substituted band without training, and ILVR [3] conditions a pretrained DDPM on a reference at a chosen downsampling factor without further learning. Prasain [33] conditions on wavelet bands instead of edge maps for map-to-satellite synthesis. None of them randomises the condition’s *bandwidth* during training so that the band becomes a control at inference. Randomising a conditioning parameter while training and exposing it at inference is not itself new: SDXL [32] samples crop coordinates during data loading and surfaces them as a user control, and Imagen [37] draws the augmentation level of its low-resolution condition at random in training and sweeps it at inference, there to find the best sample quality rather than to give the user a control. Our delta is the choice of parameter and its measured separation from adapter strength: bandwidth and adapter strength are different controls (Sec. 4.3).

Drawings as soft constraints. Sketch-based face synthesis has treated rough drawings as soft constraints [1] and, closest to our goal, lets the user set how far a drawing is trusted through a refinement level [46]; diffusion-

based stroke and sketch synthesis likewise lets the user decide the level of faithfulness to the input [2]. With pretrained diffusion models, sketch guidance can be applied at inference through a small latent guidance predictor instead of a dedicated encoder [42], and the noise added before denoising balances faithfulness to the input against realism without training [24]. An extended difference-of-Gaussians operator turns a photograph into line drawings across a range of styles set by its parameters [44]. Our control differs in what it varies: the information content of the condition, set by a frequency cutoff, on one adapter whose response we measure separately from its strength. Because the condition is a continuum and not a fixed extractor, a single adapter can expose it as a dial. StructDiff [10] targets structure-preserving, spatially controllable single-image generation, and sketch-conditioned pipelines have long used edge maps as the structural carrier.

Training and evaluation. We train with classifier-free guidance [12] dropout and Min-SNR loss weighting [9], and evaluate with FID [11] under the clean-FID protocol [29], LPIPS [51] and CLIP score [35] on FFHQ [17] with BLIP [18] captions.

3 Method

Condition. Given an image x , we convert to grayscale, take the 2D FFT, zero every frequency inside a disc of radius $r \cdot \min(H, W)/2$, invert, and min-max normalise to $[0, 1]$. The normalisation is deliberate: it removes the absolute scale of the filter response, so the model has to infer structure density from spatial statistics. We hypothesise that this is what lets a hand drawing, which carries no r , be used at test time; we do not ablate the normalisation.

Condition encoder. A convolutional encoder reads the condition image directly at 512^2 and reduces it $8\times$ to the latent grid: seven 3×3 convolutions with SiLU activations, three of them with stride 2 and followed by group normalisation, widening from 32 to 128 channels before a final projection to 4. Each attention site reads the encoder output through a 1×1 projection. The ablations call an optional 3×3 convolution of 64, 128 or 256 channels inserted before that projection a *stem*; the main model has none. We do not encode the condition with the frozen VAE; Sec. 4.4 shows why.

Injection. Encoder features enter the UNet in two places. Zero-initialised 1×1 heads produce one residual per decoder skip tensor and one for the mid block, added through the UNet’s existing ControlNet interface. In parallel, each cross-attention site attends over a 3×3 window of the condition features, gated by a zero-initialised per-channel vector. The window is not what carries the effect: at 1×1 the softmax runs over a single position, the site reduces to a gated value projection, and structure consistency is no lower (Sec. 4.5). Both paths start as exact no-ops: with them switched off the network reproduces stock Stable Diffusion 1.5 [36] bit for bit, which we assert before every training run.

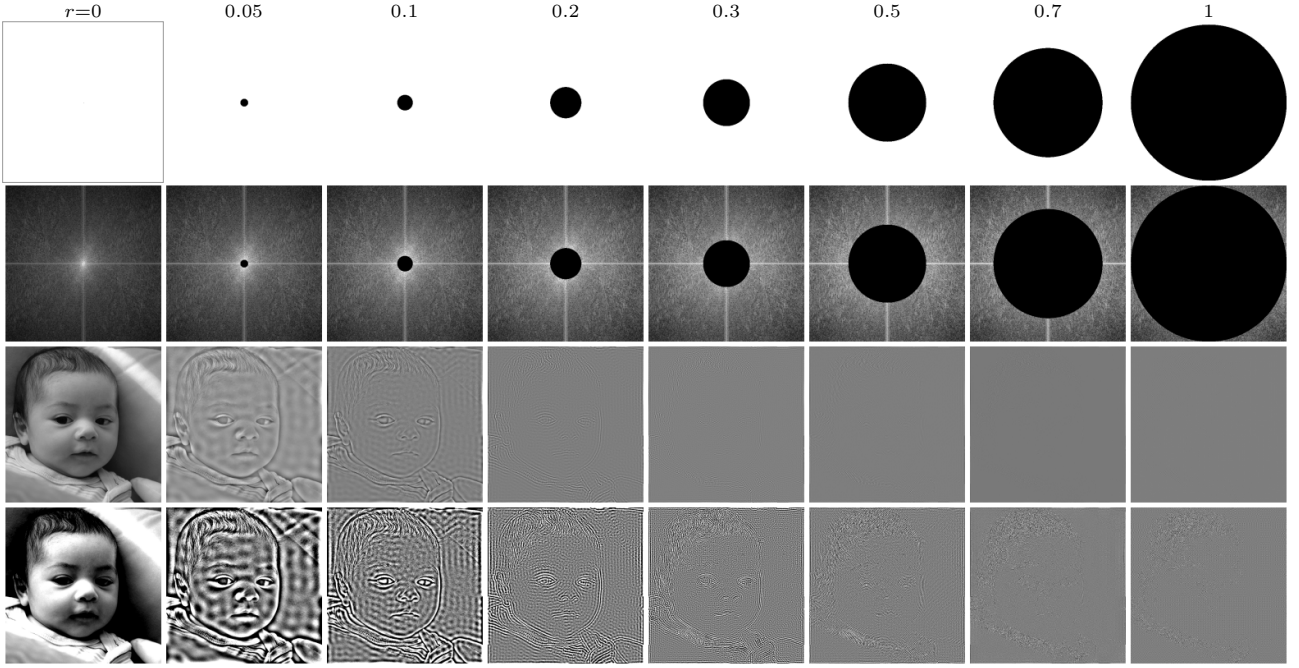


Figure 1: The dial, from the frequency side. Columns are $r = 0, 0.05, 0.1, 0.2, 0.3, 0.5, 0.7, 1.0$. **Row 1:** the radial mask; black is removed. **Row 2:** the log-magnitude spectrum after masking, with the hole growing outward. **Row 3:** the resulting condition, exactly as the model receives it. **Row 4:** the same data displayed at $\pm 2\sigma$ so that the residual content is visible. Above $r \approx 0.2$ the condition looks like flat grey (averaged over 500 test images, its standard deviation falls from 0.238 at $r=0$ to 0.034 at $r=1$), yet the face outline survives in row 4, and the model still reaches SC 0.4019 at $r=0.2$. The per-image normalisation leaves the model only this low-contrast residue to read structure from.

Training. SD1.5 [36] stays frozen. We train 6,611,469 parameters on FFHQ-512 [17] (images under their Flickr authors’ Creative Commons or public-domain licences, the dataset under CC BY-NC-SA 4.0) with BLIP [18] captions, split 65,100 train / 1,400 validation / 3,500 test by a seeded shuffle of the 70,000 images (5% test, 2% validation, seed 2026). AdamW, learning rate 10^{-4} , $\beta = (0.9, 0.999)$, cosine schedule with 500 warm-up steps, gradient clipping at 1.0, batch 32, bf16, 12 epochs on one H200. The loss is ϵ -prediction with Min-SNR weighting [9], $\gamma = 5$. Text conditioning is dropped with probability 0.1 and the structure condition independently with probability 0.05. An EMA of the adapter is kept with a 10-step interval. r is drawn uniformly from $[0, 1]$ per sample, although we claim the dial only on $r \in [0.05, 0.3]$. Code and adapter weights accompany the submission.

4 Experiments

4.1 Metrics

We evaluate on 500 of the 3,500 held-out images with a DDIM sampler at 50 steps, CFG 7.5, seed 2026.

Structure consistency. SC converts the generated image to grayscale, applies the *same* high-pass filter at the *same* cutoff the model was given, and takes the Pearson correlation against the given condition over the flattened pixel vector, per image, averaged over images. Both signals are the per-image min-max normalised filter output, so the correlation is scale free. Two anchors bound it: feeding

the ground truth back gives exactly 1.000 at every r , and an unrelated face gives 0.122 at $r=0$ and ≈ -0.02 above it. Because the instrument moves with r , a row-to-row difference mixes model behaviour with instrument change, so we report a cutoff-independent metric beside it.

Edge-F1. Canny edges (thresholds 100 and 200 on the 8-bit grayscale image) are taken from the generation and from the reference and matched with a 2-pixel tolerance, implemented as a 5×5 dilation of each map before intersecting; F1 is computed per image and averaged. It does not depend on r , which is why it appears in Table 1: 0.7796 at $r=0.05$ falling to 0.4724 at $r=1$, so the dial survives a cutoff-independent instrument, with a floor near 0.4724 that aligned FFHQ faces supply for free.

Diversity. Mean pairwise LPIPS (AlexNet backbone) over $K=4$ samples generated from one condition with different seeds; higher means the condition constrains less. It is not independent of the LPIPS column: across the sweep LPIPS – Diversity is 0.0320 ± 0.0038 , nearly constant over a range of 0.20, so the two carry close to the same information and we draw no separate conclusion from Diversity alone.

FID and CLIP. Clean-FID [11, 29] between the 500 generations and the same 500 references, and CLIP score with ViT-B/32 [35] against each image’s own caption.




Figure 2: (a) A disc of radius r is masked out of the spectrum and a trainable convolutional encoder reads the result at full resolution, never through the VAE. Its features enter the frozen UNet twice: as zero-initialised residuals on every skip tensor ($\oplus$), and through StandAlone attention at the 16 cross-attention sites. Stroke weight follows the measured split: at $r=0$, SC is 0.5370 with only the skip path active and 0.2112 with only attention (Sec. 4.5). (b) One StandAlone site. The query comes from the UNet’s own latent through the frozen W_Q ; keys and values come from the condition through trained 1×1 projections, unfolded so each position attends over its 3×3 neighbourhood rather than the whole map; a 1×1 window performs at least as well (Table 3). The result passes the frozen output projection and a zero-initialised per-channel gate, then is added both into the latent — so the text query below sees structure — and to what the layer returns.

4.2 The dial

Table 1 gives the operating curve. On $r \in [0.05, 0.3]$ structure consistency falls at every step and sample diversity rises at every step, which is the trade-off the dial is meant to expose. Two qualifications belong with that sentence. Going from $r=0$ to $r=0.05$ structure consistency *rises* by 0.0479 ($z = 6.24$, paired within one checkpoint), so the sweep is not monotone at its left end; the same rise appears in Edge-F1 (0.7776 to 0.7796), so it is model behaviour rather than an artefact of the SC instrument. And at the right end $r=0.5$ and $r=0.7$ are not separable across seeds (paired $t = 0.732$, $df = 2$, critical 4.303). FID stays between 50.13 and 57.40 across the entire sweep.

Figure 3 holds the seed and the caption fixed and varies only r . Because the caption is derived from the reference, it could be carrying part of the structure. We repeated the sweep with one generic caption, “a photo of a person”, for every image. Structure consistency barely moves — 0.7039, 0.6179 and 0.2196 at $r=0$, 0.1 and 0.3, against 0.6991, 0.6211 and 0.2347 with captions — while FID rises to 55.56–65.84, as it should once the caption stops describing the image. The structure comes from the condition. Figure 4 plots the curve and, below it, both dials in the same structure-consistency/FID plane.

Across seeds 2026, 2027 and 2028 the curve is stable: SC at $r=0.05$ is 0.7479 ± 0.0015 and at $r=0.3$ it is 0.2318 ± 0.0106 .

4.3 Strength is not the same knob

An obvious alternative is to train on one cutoff and turn the adapter’s contribution down at inference, the knob lightweight control modules expose [6]. We ran it on the adapter trained at fixed $r=0.1$ (Sec. 4.6), scaling

r	SC $\pm$ SEM	seed sd	Edge-F1	Diversity	FID
0.00	0.6991 ± 0.0070	0.0105	0.7776	0.2994	50.13
0.05	0.7469 ± 0.0030	0.0015	0.7796	0.3126	51.58
0.10	0.6211 ± 0.0035	0.0059	0.7439	0.3414	54.68
0.20	0.4019 ± 0.0037	0.0098	0.7113	0.3719	53.45
0.30	0.2347 ± 0.0041	0.0106	0.6878	0.3794	52.77
<i>beyond the range we claim</i>					
0.50	0.1259 ± 0.0047	0.0093	0.6225	0.4100	53.97
0.70	0.1110 ± 0.0047	0.0186	0.5788	0.4334	53.69
1.00	0.0010 ± 0.0006	0.0002	0.4724	0.5066	57.40

Table 1: Operating curve, 500 of the 3,500 test images. SEM is over those 500 images; seed sd is over three training seeds. Edge-F1 is a cutoff-independent reading of the same runs and is reported alongside SC because SC re-extracts at the same cutoff the model was given (Sec. 4.1). FID has a single estimate per cell and no variance, so small differences between rows are not read as effects.

both injection paths, and, as a second check, scaled the randomised- r model at the same cutoff.

Both scaled curves leave the image distribution as they go. At scale 0.25 the fixed- r adapter reaches FID 102.43 and diversity 0.7689, beyond anything the cutoff produces over its whole sweep, where FID stays within 50.13–57.40 and diversity peaks at 0.5066. Turning the adapter down removes structure by degrading the sample; narrowing the condition’s band removes it while image quality holds.

4.4 The VAE eats the condition

Encoding the condition with the frozen SD VAE loses most of it. An image-prompt design such as IP-Adapter [47] takes a different route — a CLIP image encoder and a projection — so this measurement does not speak to it.

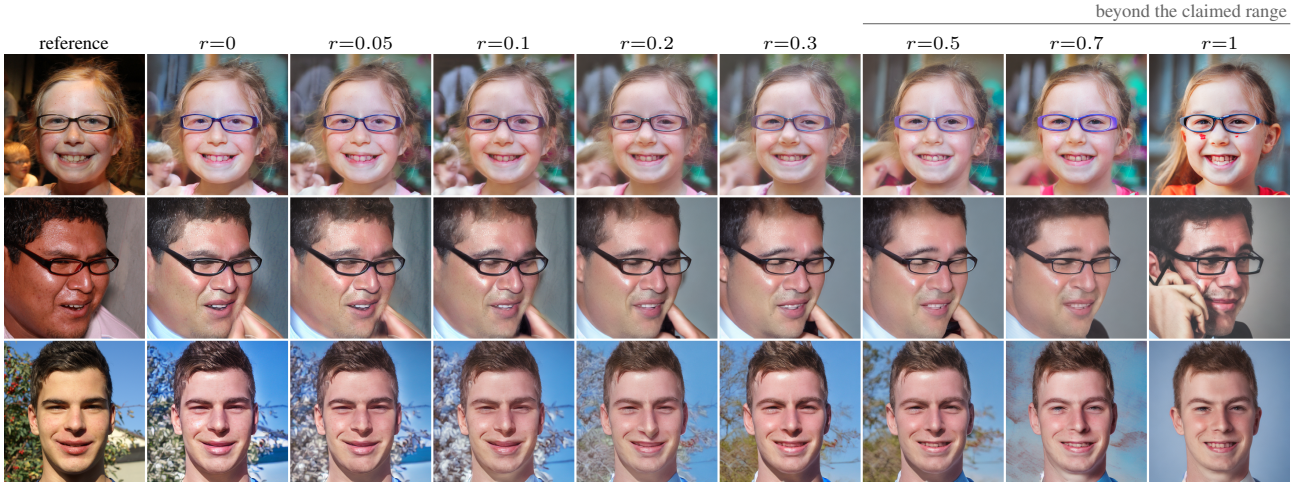


Figure 3: One seed and one caption per row; only r changes. The rows are the first three test images in generation order. As r rises, background detail thins and clothing and colour drift from the reference while the face layout holds; at $r=1$ the pose of the second row changes as well. The reference is never shown to the model. Columns beyond $r=0.3$ lie outside the claimed range; at $r=1$ colour artefacts appear on the first row.




Figure 4: **Top:** structure consistency falls and sample diversity rises as the cutoff grows. **Bottom:** the cutoff sweep of Table 1 and the two scaled adapters of Table 2 in the structure-consistency / FID plane. Narrowing the condition’s bandwidth traverses the axis at nearly constant FID; weakening the adapter buys the same structure consistency by degrading the image. Both panels are plotted directly from the evaluation outputs.

Figure 5 plots it. Passing the condition through encode-decode and re-measuring leaves 0.8562 at $r=0.1$, 0.5773 at $r=0.3$ and 0.0700 at $r=0.7$. The round trip lower-bounds what the latent retains — a component may survive there that the frozen decoder cannot render — but an adapter reading that latent cannot recover what the round trip has lost. Building the condition at 256 and upsampling moves the structure to lower absolute frequency and preserves far more (0.7978 at $r=0.3$), which is a data-pipeline artefact worth knowing about when comparing against older set-ups.

Setting	SC	Diversity	FID
<i>fixed-r=0.1 adapter, scaled</i>			
scale 0.25	0.0640	0.7689	102.43
scale 0.50	0.3193	0.5540	70.91
scale 0.75	0.5829	0.4044	57.52
scale 1.00	0.7050	0.2967	48.51
<i>randomised-r adapter at $r=0.1$, scaled</i>			
scale 0.25	0.0408	0.7752	104.95
scale 0.50	0.2146	0.6061	76.46
scale 0.75	0.4893	0.4171	58.90
<i>randomised-r adapter, cutoff moved</i>			
$r=0.2$	0.4019	0.3719	53.45
$r=0.3$	0.2347	0.3794	52.77

Table 2: Strength against cutoff. The fixed- r adapter at scale 0.50 reaches SC 0.3193 at FID 70.91; moving the cutoff brackets that value with $r=0.2$ (0.4019, FID 53.45) and $r=0.3$ (0.2347, FID 52.77), both at lower FID. Scaling the randomised model behaves the same way.

4.5 Where to inject

Table 4 lists absolute configurations; each row below the control differs from the *conv / attn* row in one field. Skip-connection residuals dominate. Figure 6 decomposes them. Ablating them away and keeping only cross-attention drops SC at $r=0$ from 0.6991 to 0.2112; keeping only the skip path retains 0.5370. With both paths off the model collapses to 0.0368 SC and FID 120.50, which is the control that says the adapter is doing the work. The split is not the same on both axes: on SC the skip path leads by roughly $2.5\times$, but on FID the two paths are almost indistinguishable (65.12 against 67.07), so the placement effect we measure is a structure-fidelity effect and not a general image-quality one.

Capacity helps much less: widening the projection to a 128-channel stem raises SC at $r=0.05$ from 0.1485 to 0.2826, and doubling again to 256 reaches only 0.2986. A 7×7 window gives 0.0359, far below the 3×3 design




Figure 5: Correlation between the condition fed to the autoencoder and the same condition re-extracted after an encode-decode round trip, 24 images. Above $r \approx 0.5$ the condition is effectively gone, so an adapter reading the VAE latent cannot recover it. Building the condition at 256 and upsampling puts the structure at lower absolute frequency and preserves far more, which is why comparisons against older pipelines need the source resolution stated.




Figure 6: Both injection paths on, then each alone, then neither. The skip residuals carry most of the structure signal; cross-attention alone recovers little of it, and with both off the model falls to the frozen backbone (FID 120.50 against 50.13 with both paths on; skip only 65.12, attention only 67.07). The dashed line is the both-off control at $r=0$. At $r=0$ the single-path gains over the control roughly add up to the both-paths score; at $r=0.3$ both paths together exceed that sum.

(0.1485) and close to the no-adapter control (0.0074). All of these vary capacity with the skip path absent, so they bound how far capacity carries *that* design rather than settling capacity against placement in general.

Table 3 repeats the one-field ablations on the working design with both paths on. Each is a single seed, so we read it against the main model’s three-seed spread of about 0.01 in SC. Capacity helps here too, modestly: a 256-channel stem lifts SC at $r=0.05$ from 0.7469 to 0.7631 and lowers FID at $r=0.1$ from 54.68 to 51.36, small next to adding the skip path (0.2112 to 0.6991 at $r=0$). The attention window does not help: 1×1 matches 3×3 within the seed spread or exceeds it at every cutoff listed, and 5×5 is lower at all of them, so what the attention sites contribute is where the condition enters, not a neighbourhood. The zero-initialised gate is worth keeping; without it SC is lower from $r=0$ to 0.3 and FID at $r=0.1$ is higher. The largest single change is how r is drawn. A log-uniform sampler lifts SC at $r=0$ to 0.9480 and lowers FID there to 42.06, but gives up the upper range (0.0858 at $r=0.5$ against 0.1259): it moves the dial’s resolution toward small r rather than improving it everywhere. The results above use the uniform sampler. The main model keeps the 3×3 window and no stem, although the 1×1 and stem variants score higher at several cutoffs in Table 6, because these one-field ablations were

One field changed	SC				FID
	$r=0$	0.05	0.3	0.5	$r=0.1$
main model	0.6991	0.7469	0.2347	0.1259	54.68
seed sd (3 seeds)	0.0105	0.0015	0.0106	0.0093	–
window 1×1	0.7825	0.7640	0.2271	0.1255	55.70
window 5×5	0.6764	0.7175	0.2039	0.0944	55.68
no gate	0.6838	0.7322	0.2058	0.1290	57.02
avg-pool downsampling	0.7203	0.7465	0.2279	0.1241	55.13
stem 64	0.7775	0.7594	0.2354	0.1170	51.91
stem 256	0.7900	0.7631	0.2511	0.1214	51.36
colour condition	0.6949	0.7463	0.2343	0.1255	54.75
no condition dropout	0.7042	0.7353	0.2279	0.1231	54.69
$r = u^2$	0.8936	0.7836	0.2468	0.1039	53.46
r log-uniform	0.9480	0.8000	0.2302	0.0858	54.87

Table 3: One-field ablations on the working model, skip and attention paths both on, one training seed each; the second row is the main model’s spread over three seeds. The main model uses a 3×3 window, a zero-initialised gate, learned 2×2 strided downsampling of the condition features at each attention site, no stem before the 1×1 projection, a grayscale condition, condition dropout 0.05, and r uniform on $[0, 1]$. Drawing $r = u^2$ with u uniform concentrates training at small r ; the log-uniform sampler on $[10^{-3}, 1]$ does so more strongly.

Cond. encoder	Injection	SC $r=0.05$	SC $r=0.1$	FID $r=0.1$
no adapter (control)		0.0074	0.0082	120.50
frozen VAE	attn, no gate	0.0234	0.0147	76.25
frozen VAE	attn	0.0303	0.0177	65.71
conv	attn	0.1485	0.0772	56.39
conv, 128	attn	0.2826	0.1882	53.17
conv, 256	attn	0.2986	0.2158	52.84
conv	attn, 7×7	0.0359	0.0218	58.51
conv	attn + skip	0.7469	0.6211	54.68

Table 4: Absolute configurations, not a cumulative chain: every row below the control differs from the *conv* / *attn* row in exactly one field, and all rows use a 3×3 window and a zero-initialised gate unless stated. The main model is the last row and carries no stem; its attention sites read the encoder output through a 1×1 projection. The frozen-VAE rows use attention injection only. They and the 7×7 row stay close to the no-adapter control at both cutoffs, so at that level a weak design and a run that failed to learn are not distinguishable. Stem width with the skip path present is varied in Table 3.

trained after the main configuration and its three seeds had been fixed.

4.6 Against per-cutoff specialists

We trained four adapters at fixed r and compared each against the single randomised- r model at its own cutoff. Because the single model was trained three times, we express each gap against its three-seed spread rather than against a within-run standard error; a gap smaller than the two-sided t critical value at $df = 2$ (4.303) is not separable. The result is mixed: the specialist wins at $r=0.1$ (0.7049 against 0.6245, $\Delta = +0.0803$, +13.55 sd) and the single model wins at $r=0.3$ (0.2318 against 0.1214, -0.1104 , -10.42 sd). Neither $r=0.5$ (-1.00 sd) nor $r=0.7$ ($+0.14$ sd) separates.

One model does not dominate four. It gives a dial the

r	skip only	both paths	attention only
0.05	0.7085	0.7469	—
0.10	0.5765	0.6211	0.0275
0.20	0.3508	0.4019	—
0.30	0.1608	0.2347	0.0279
0.50	0.0709	0.1259	—
0.70	0.0757	0.1110	—

Table 5: Structure consistency for a model trained with the skip path only, against the working model and against its cross-attention branch measured in isolation. The last column comes from switching a path off at test time, so it is not a trained-from-scratch comparison.

specialists cannot: a model trained at $r=0.1$ has never seen $r=0.3$, so its curve is a single point.

4.7 The skip path alone

Section 4.5 ablates the two injection paths on a model trained with both. That confounds placement with capacity: the cross-attention branch is still present and trained, just switched off at test time. We therefore trained the skip path from scratch, with the attention branch removed entirely.

It reaches SC 0.5765 at $r=0.1$ against 0.6211 for both paths and 0.0275 for cross-attention alone, at FID 58.20 against 54.68. A single path trained on its own therefore recovers most, but not all, of the combined model: the paths are complementary rather than redundant, and placement still dominates at matched capacity. The full curve is in Table 5.

4.8 Reading the dial with a fixed instrument

Structure consistency re-extracts the condition at the same cutoff the model was given, so the measuring instrument moves with the thing it measures. A falling curve could in principle come from the instrument rather than the model. We therefore re-scored the same generations with the cutoff held at $r_m=0.1$ throughout.

The curve still falls: 0.6524 at $r=0.05$, 0.6210 at $r=0.1$ and 0.2599 at $r=0.3$. Holding the instrument fixed compresses the range, as it must, but does not remove the decline, so the dial is a property of the generations and not of the metric.

4.9 Outside the training domain

Everything above is FFHQ. To test whether the dial is specific to faces we ran the same sweep on photographs from outside the training distribution, scoring against the input drawing with the identical script, prompt and cutoffs.

Structure consistency falls from 0.4506 at $r=0.05$ to 0.1384 at $r=0.3$, a decline of 0.3122 against 0.1839 in domain — a ratio of 1.698. The dial therefore carries out of domain, and if anything responds more sharply there, which is what one would expect when the model has less of a prior to fall back on. With the adapter off the same images score 0.0145, so the residual structure is coming from the condition and not from the backbone.



Figure 7: Hand-drawn face sketches (top row) and what the model makes of them (bottom row), zero-shot. The drawings carry no cutoff; the model has to read structure density off the drawing itself, which the per-image min-max normalisation permits. Most of the scans carry the archive filename stamped across the top, and the model redraws it — a literal reading of the condition it was given. Removing the whole top 20% of the frame (102 rows), a generous over-estimate of that band, moves the score to 0.5398 (-0.0086, 0.32 standard errors), so the band is not what the score is measuring.

Image quality does not transfer: FID rises from 65.71 in domain to 148.97 out of it. The control survives the domain shift; the generative quality does not, and we make no claim otherwise.

4.10 Baselines

Against pretrained ControlNet [50] on its own condition, over the same test images and captions: ControlNet-Canny reaches Edge-F1 0.6854, LPIPS 0.5294, FID 71.92 with 361,279,120 trainable parameters; ControlNet-HED reaches 0.5733 / 0.6548 / 118.98. At $r=0.1$ our adapter reaches Edge-F1 0.7439, LPIPS 0.3748 and FID 54.68 with 6,611,469 parameters. Two asymmetries make this a comparison of operating points rather than a like-for-like control: the conditions differ (Canny edges against a frequency band), and our adapter was trained on 65,100 FFHQ images while ControlNet is zero-shot here and FID is measured against FFHQ statistics. On text alignment we do not win anywhere: our best CLIP over the whole sweep is 30.30, below ControlNet-Canny’s 30.34 (ControlNet-HED reaches 29.16). T2I-Adapter-Canny [25], a 77,000,640-parameter side network and so closer to our design than a copied encoder, reaches Edge-F1 0.5845, LPIPS 0.5690, FID 86.66 and CLIP 30.58 under the same protocol. We leave out its sketch checkpoint, which expects a PiDiNet sketch we did not compute; fed Canny edges it would measure that mismatch rather than the method. Training ControlNet on our condition and data would be the matched control; we have not run it.

4.11 Hand-drawn sketches

Figure 7 shows the transfer. On 40 real face drawings from the CUHK Face Sketch database [43], which the model was never trained on, prompted with “a high-quality photo of a face”, structure consistency against the input drawing is 0.5484 and CLIP score is 23.43. FID is not reported: 40 images is too few for a stable estimate, and our evaluator

refuses rather than print one. We also ran 500 general-scene drawings, the sketch half of the public photograph-sketch set `rhfeiyang/photo-sketch-pair-500` on Hugging Face, far outside the training domain and with the same face prompt, where structure consistency falls to 0.1714 and FID against FFHQ is 285.91: the model follows the drawing but does not produce a face. That is the expected failure for a face model and we report it rather than omit it.

A drawing arrives without a cutoff, but the artist can give it one: filter the drawing at r and hand the model the result, the operation every training condition went through. On 40 face drawings, structure consistency against the original drawing falls from 0.5486 at $r=0$ to 0.3390 at $r=0.1$, 0.2271 at $r=0.3$ and 0.0407 at $r=1$, so the dial carries over to drawings; $r=0.5$ and $r=0.7$ lie within a standard error (0.0419) of each other. Against the filtered drawing it was actually given, the score at $r=0.1$ is 0.1399, well below the 0.6211 a filtered photograph yields; we have not isolated why.

5 Limitations

The useful range is roughly $r \in [0, 0.3]$. Above it the condition retains little energy — 0.0016 of the spectrum at $r=0.5$ — and structure consistency becomes hard to interpret, since two nearly empty signals correlate near zero whatever the model does.

Everything here is one domain (FFHQ faces) and one backbone (SD1.5), and the three results travel differently. The dial is a property of the condition and need not depend on the backbone, but we have not tested another. The skip-path injection does not transfer as it stands: MMDiT backbones [4] have no UNet skip connections, and control there enters through shared tokens or LoRA [14] modules [22, 40, 52]. The VAE measurement was made on the four-channel SD autoencoder only. Our adapter is also not small by current standards: Simple-ControlNet [6] reports under 3M parameters on SD1.5, ControlNet-XS [49] ships much smaller control networks than ControlNet, and the DiT modules above use a fraction of our count.

The one-field ablations of Table 3 are single-seed. The specialist comparison is honest but unflattering at $r=0.1$, and we have not run a user study.

6 Conclusion

Randomising the bandwidth of a structure condition during training turns the cutoff into an inference-time control, and that control is not reproducible by scaling the adapter down. Getting there required fixing where the condition enters the network and keeping it out of the frozen autoencoder; both were measured on the SD1.5 autoencoder and UNet, and we expect them to carry to conditioning designs built on that family. We have not tested others.

A Full per-cutoff results

The main text reports each experiment at the cutoff that carries its argument. The tables here give every cutoff that was measured, from the same `results.json` files the rest of the paper is generated from.

variant	$r=0$	0.05	0.1	0.3	0.5
working model	0.6991	0.7469	0.6211	0.2347	0.1259
1×1 window	0.7825	0.7640	0.6382	0.2271	0.1255
5×5 window	0.6764	0.7175	0.5919	0.2039	0.0944
64-ch stem	0.7775	0.7594	0.6358	0.2354	0.1170
256-ch stem	0.7900	0.7631	0.6332	0.2511	0.1214
average downsample	0.7203	0.7465	0.6181	0.2279	0.1241
colour condition	0.6949	0.7463	0.6251	0.2343	0.1255
log cutoff sampler	0.9480	0.8000	0.6617	0.2302	0.0858
sqrt cutoff sampler	0.8936	0.7836	0.6568	0.2468	0.1039
no condition dropout	0.7042	0.7353	0.6108	0.2279	0.1231
no zero-init gate	0.6838	0.7322	0.6118	0.2058	0.1290
<i>designs not adopted (Sec. 4.4)</i>					
256px source pipeline	—	0.0271	0.0171	0.0291	—
VAE condition	—	0.0234	0.0147	0.0172	—

Table 6: Structure consistency for every one-field ablation, at every measured cutoff. Each row changes one field of the working model and nothing else. The last block holds two designs the paper argues against: building the condition at 256 and upsampling (FID 86.84 at $r=0.1$), and reading it through the frozen VAE.

variant	$r=0$	0.05	0.1	0.3	0.5
working model	50.13	51.58	54.68	52.77	53.97
1×1 window	49.13	52.34	55.70	54.16	54.33
5×5 window	52.05	54.16	55.68	54.43	55.39
64-ch stem	48.92	50.25	51.91	51.08	52.03
256-ch stem	47.74	49.49	51.36	51.77	52.13
average downsample	50.90	53.43	55.13	54.29	54.32
colour condition	52.53	52.09	54.75	52.67	52.84
log cutoff sampler	42.06	52.17	54.87	54.33	56.96
sqrt cutoff sampler	44.69	51.89	53.46	52.87	54.52
no condition dropout	51.17	52.50	54.69	53.13	54.26
no zero-init gate	52.72	55.63	57.02	54.53	55.51

Table 7: FID for the same one-field ablations as Table 6, one training seed each. FID is a single estimate per cell with no variance (Table 1), so small differences are not read as effects; the largest departures are the log-uniform and $r=u^2$ samplers at $r=0$ (42.06 and 44.69 against 50.13).

References

- [1] Shu-Yu Chen, Wanchao Su, Lin Gao, Shihong Xia, and Hongbo Fu. Deepfacedrawing: Deep generation of face images from sketches. *ACM Transactions on Graphics*, 39(4), 2020. doi: 10.1145/3386569.3392386.
- [2] Shin-I Cheng, Yu-Jie Chen, Wei-Chen Chiu, Hung-Yu Tseng, and Hsin-Ying Lee. Adaptively-realistic image generation from stroke and sketch with diffusion model. In *IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, pages 4054–4062, 2023. doi: 10.1109/WACV56688.2023.00404.

r	SC	FID	LPIPS	CLIP	Edge-F1	Div.
0.00	0.6991	50.13	0.3341	29.54	0.7776	0.2994
0.05	0.7469	51.58	0.3495	29.41	0.7796	0.3126
0.10	0.6211	54.68	0.3748	29.53	0.7439	0.3414
0.20	0.4019	53.45	0.4003	29.47	0.7113	0.3719
0.30	0.2347	52.77	0.4148	29.73	0.6878	0.3794
0.50	0.1259	53.97	0.4418	29.78	0.6225	0.4100
0.70	0.1110	53.69	0.4625	29.92	0.5788	0.4334
1.00	0.0010	57.40	0.5329	30.30	0.4724	0.5066

Table 8: Every metric for the working model, at every cutoff. The SC and FID columns are the ones plotted in Figure 4.

r	skip only	skip only FID	seed mean	seed s.d.
0.00	0.6284	55.84	0.7109	0.0105
0.05	0.7085	57.65	0.7479	0.0015
0.10	0.5765	58.20	0.6245	0.0059
0.20	0.3508	57.07	0.4036	0.0098
0.30	0.1608	57.21	0.2318	0.0106
0.50	0.0709	57.41	0.1304	0.0093
0.70	0.0757	57.95	0.1261	0.0186
1.00	0.0006	59.60	0.0012	0.0002

Table 9: Left: the skip path trained alone (Sec. 4.7). Right: mean and standard deviation of structure consistency over seeds 2026, 2027 and 2028 for the working model.

adapter strength	SC	FID	SC, seed 2027
0.1	0.0165	117.58	—
0.2	0.0300	110.20	—
0.25	0.0408	104.95	0.0309
0.35	0.0808	95.62	—
0.5	0.2146	76.46	0.2038
0.6	0.3373	66.39	—
0.75	0.4893	58.90	0.4823
1.00	0.6211	54.68	—

Table 10: The strength knob on the randomised- r adapter at $r=0.1$, swept finely. Each scaled point below 1.00 pays for lower structure consistency with FID; over $r \in [0.05, 0.3]$ the cutoff dial keeps FID between 51.58 and 54.68 (Table 8). Scales 0.25, 0.5 and 0.75 were repeated with sampling seed 2027 on the same checkpoint.

r	specialist	one model, 3 seeds		difference	diff./sd
		mean	sd		
0.10	0.7049	0.6245	0.0059	+0.0803	+13.55
0.30	0.1214	0.2318	0.0106	-0.1104	-10.42
0.50	0.1211	0.1304	0.0093	-0.0093	-1.00
0.70	0.1287	0.1261	0.0186	+0.0026	+0.14

Table 11: Per-cutoff specialists (one training run each, seed 2026) against the single randomised- r model averaged over training seeds 2026–2028, each at its own cutoff. Difference is specialist minus the three-seed mean; the last column divides it by the three-seed standard deviation, as in Sec. 4.6. A positive difference favours the specialist. The specialist wins at $r=0.1$, the single model wins at $r=0.3$, and neither separates at $r=0.5$ or 0.7 ; what the single model gives that no specialist can is the rest of the curve.

- [3] Jooyoung Choi, Sungwon Kim, Yonghyun Jeong, Youngjune Gwon, and Sungroh Yoon. Ilvr: Conditioning method for denoising diffusion probabilistic models. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 14367–14376, 2021. doi: 10.1109/ICCV48922.2021.01410.
- [4] Patrick Esser, Sumith Kulal, Andreas Blattmann, Rahim Entezari, Jonas Müller, Harry Saini, Yam Levi, Dominik Lorenz, Axel Sauer, Frederic Boesel, Dustin Podell, Tim Dockhorn, Zion English, Kyle Lacey, Alex Goodwin, Yannik Marek, and Robin Rombach. Scaling rectified flow transformers for high-resolution image synthesis. *arXiv:2403.03206*, 2024. URL <https://arxiv.org/abs/2403.03206>.
- [5] Zeev Farbman, Raanan Fattal, Dani Lischinski, and Richard Szeliski. Edge-preserving decompositions for multi-scale tone and detail manipulation. *ACM Transactions on Graphics*, 27(3), 2008. doi: 10.1145/1360612.1360666.
- [6] Zhida Feng, Li Chen, Yuenan Sun, Jiaxiang Liu, and Shikun Feng. Simplifying control mechanism in text-to-image diffusion models. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 3013–3021, 2025. doi: 10.1609/aaai.v39i3.32309.
- [7] Xiang Gao and Jiaying Liu. Fbsdiff: Plug-and-play frequency band substitution of diffusion features for highly controllable text-driven image translation. In *Proceedings of the 32nd ACM International Conference on Multimedia*, pages 4101–4109, 2024. doi: 10.1145/3664647.3681380.
- [8] Xiang Gao, Zhengbo Xu, Junhan Zhao, and Jiaying Liu. Frequency-controlled diffusion model for versatile text-guided image-to-image translation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 1824–1832, 2024. doi: 10.1609/aaai.v38i3.27951.
- [9] Tiankai Hang, Shuyang Gu, Chen Li, Jianmin Bao, Dong Chen, Han Hu, Xin Geng, and Baining Guo. Efficient diffusion training via min-snr weighting strategy. *arXiv:2303.09556*, 2023. URL <https://arxiv.org/abs/2303.09556>.
- [10] Yinxi He, Kang Liao, Chunyu Lin, Tianyi Wei, and Yao Zhao. Structdiff: A structure-preserving and spatially controllable diffusion model for single-image generation. *arXiv:2604.12575*, 2026. URL <https://arxiv.org/abs/2604.12575>.
- [11] Martin Heusel, Hubert Ramsauer, Thomas Unterthiner, Bernhard Nessler, and Sepp Hochreiter. Gans trained by a two time-scale update rule converge to a local nash equilibrium. *arXiv:1706.08500*, 2017. URL <https://arxiv.org/abs/1706.08500>.
- [12] Jonathan Ho and Tim Salimans. Classifier-free diffusion guidance. *arXiv:2207.12598*, 2022. URL <https://arxiv.org/abs/2207.12598>.

r	OOD SC	OOD FID	in-domain SC	in-domain FID
0.05	0.4506	128.40	0.5981	61.07
0.10	0.3106	148.97	0.5057	65.71
0.20	0.1944	154.24	0.4407	68.78
0.30	0.1384	160.82	0.4142	70.75

Table 12: Out-of-domain photographs against in-domain, scored against the input drawing with the same script and prompt (Sec. 4.9). The control falls in both; the image quality does not transfer.

r	sketch dial SC	$\pm$ s.e.m.	specialist
0.00	0.5486	0.0273	—
0.05	0.4826	0.0295	—
0.10	0.3390	0.0356	0.7049
0.20	0.2746	0.0360	—
0.30	0.2271	0.0364	0.1214
0.50	0.1379	0.0419	0.1211
0.70	0.1232	0.0419	0.1287
1.00	0.0407	0.0326	—

Table 13: Left: the dial applied to 40 hand-drawn sketches, scored against the drawing. Right: per-cutoff specialists at their own cutoff, for the four cutoffs where one was trained.

- [13] Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. *arXiv:2006.11239*, 2020. URL <https://arxiv.org/abs/2006.11239>.
- [14] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models. *arXiv:2106.09685*, 2021. URL <https://arxiv.org/abs/2106.09685>.
- [15] Chao Huang, Susan Liang, Yunlong Tang, Jing Bi, Li Ma, Yapeng Tian, and Chenliang Xu. Fresca: Scaling in frequency space enhances diffusion models. *arXiv:2504.02154*, 2025. URL <https://arxiv.org/abs/2504.02154>.
- [16] Sida Huang, Siqi Huang, Ping Luo, and Hongyuan Zhang. Laytrol: Preserving pretrained knowledge in layout control for multimodal diffusion transformers. *arXiv:2511.07934*, 2025. URL <https://arxiv.org/abs/2511.07934>.
- [17] Tero Karras, Samuli Laine, and Timo Aila. A style-based generator architecture for generative adversarial networks. *arXiv:1812.04948*, 2018. URL <https://arxiv.org/abs/1812.04948>.
- [18] Junnan Li, Dongxu Li, Caiming Xiong, and Steven Hoi. Blip: Bootstrapping language-image pre-training for unified vision-language understanding and generation. *arXiv:2201.12086*, 2022. URL <https://arxiv.org/abs/2201.12086>.
- [19] Ming Li, Taojiannan Yang, Huafeng Kuang, Jie Wu, Zhaoning Wang, Xuefeng Xiao, and Chen Chen. Controlnet++: Improving conditional controls with efficient consistency feedback. *arXiv:2404.07987*, 2024. URL <https://arxiv.org/abs/2404.07987>.
- [20] Yuheng Li, Haotian Liu, Qingyang Wu, Fangzhou Mu, Jianwei Yang, Jianfeng Gao, Chunyuan Li, and Yong Jae Lee. Gligen: Open-set grounded text-to-image generation. *arXiv:2301.07093*, 2023. URL <https://arxiv.org/abs/2301.07093>.
- [21] Xin Lin, Haodong Li, Zhifei Zhang, Yutong Yang, Haitian Zheng, Juanxi Tian, Zhe Lin, and Truong Nguyen. Pixelcontrol: Fine-grained condition fidelity in text-to-image diffusion. *arXiv:2608.15705*, 2026. URL <https://arxiv.org/abs/2608.15705>.
- [22] Shanyuan Liu, Jian Zhu, Junda Lu, Yue Gong, Li-uzhuozheng Li, Bo Cheng, Yuhang Ma, Liebucha Wu, Xiaoyu Wu, Dawei Leng, and Yuhui Yin. Nanocontrol: A lightweight framework for precise and efficient control in diffusion transformer. *arXiv:2508.10424*, 2025. URL <https://arxiv.org/abs/2508.10424>.
- [23] Zehong Ma, Longhui Wei, Shuai Wang, Shiliang Zhang, and Qi Tian. Deco: Frequency-decoupled pixel diffusion for end-to-end image generation. *arXiv:2511.19365*, 2025. URL <https://arxiv.org/abs/2511.19365>.
- [24] Chenlin Meng, Yutong He, Yang Song, Jiaming Song, Jiajun Wu, Jun-Yan Zhu, and Stefano Ermon. Sdedit: Guided image synthesis and editing with stochastic differential equations. In *International Conference on Learning Representations*, 2022. URL <https://arxiv.org/abs/2108.01073>.
- [25] Chong Mou, Xintao Wang, Liangbin Xie, Yanze Wu, Jian Zhang, Zhongang Qi, Ying Shan, and Xiao-hu Qie. T2i-adapter: Learning adapters to dig out more controllable ability for text-to-image diffusion models. *arXiv:2302.08453*, 2023. URL <https://arxiv.org/abs/2302.08453>.
- [26] Alex Nichol, Prafulla Dhariwal, Aditya Ramesh, Pranav Shyam, Pamela Mishkin, Bob McGrew, Ilya Sutskever, and Mark Chen. Glide: Towards photorealistic image generation and editing with text-guided diffusion models. *arXiv:2112.10741*, 2021. URL <https://arxiv.org/abs/2112.10741>.
- [27] Aude Oliva, Antonio Torralba, and Philippe G. Schyns. Hybrid images. *ACM Transactions on Graphics*, 25(3):527–532, 2006. doi: 10.1145/1179352.1141919.
- [28] Sylvain Paris, Samuel W. Hasinoff, and Jan Kautz. Local laplacian filters: Edge-aware image processing with a laplacian pyramid. *ACM Transactions on Graphics*, 30(4), 2011. doi: 10.1145/2010324.1964963.
- [29] Gaurav Parmar, Richard Zhang, and Jun-Yan Zhu. On aliased resizing and surprising subtleties in gan evaluation. *arXiv:2104.11222*, 2021. URL <https://arxiv.org/abs/2104.11222>.

- [30] William Peebles and Saining Xie. Scalable diffusion models with transformers. *arXiv:2212.09748*, 2022. URL <https://arxiv.org/abs/2212.09748>.
- [31] Bohao Peng, Jian Wang, Yuechen Zhang, Wenbo Li, Ming-Chang Yang, and Jiaya Jia. Controlnext: Powerful and efficient control for image and video generation. *arXiv:2408.06070*, 2024. URL <https://arxiv.org/abs/2408.06070>.
- [32] Dustin Podell, Zion English, Kyle Lacey, Andreas Blattmann, Tim Dockhorn, Jonas Müller, Joe Penna, and Robin Rombach. Sdxl: Improving latent diffusion models for high-resolution image synthesis. *arXiv:2307.01952*, 2023. URL <https://arxiv.org/abs/2307.01952>.
- [33] Arisha Prasain. Beyond edge maps: Wavelet-domain conditioning for multi-adapter map-to-satellite diffusion. *arXiv:2608.00083*, 2026. URL <https://arxiv.org/abs/2608.00083>.
- [34] Can Qin, Shu Zhang, Ning Yu, Yihao Feng, Xinyi Yang, Yingbo Zhou, Huan Wang, Juan Carlos Niebles, Caiming Xiong, Silvio Savarese, Stefano Ermon, Yun Fu, and Ran Xu. Unicontrol: A unified diffusion model for controllable visual generation in the wild. *arXiv:2305.11147*, 2023. URL <https://arxiv.org/abs/2305.11147>.
- [35] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. *arXiv:2103.00020*, 2021. URL <https://arxiv.org/abs/2103.00020>.
- [36] Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models. *arXiv:2112.10752*, 2021. URL <https://arxiv.org/abs/2112.10752>.
- [37] Chitwan Saharia, William Chan, Saurabh Saxena, Lala Li, Jay Whang, Emily Denton, Seyed Kamyar Seyed Ghasemipour, Burcu Karagol Ayan, S. Sara Mahdavi, Rapha Gontijo Lopes, Tim Salimans, Jonathan Ho, David J Fleet, and Mohammad Norouzi. Photorealistic text-to-image diffusion models with deep language understanding. *arXiv:2205.11487*, 2022. URL <https://arxiv.org/abs/2205.11487>.
- [38] Chenyang Si, Ziqi Huang, Yuming Jiang, and Ziwei Liu. Freeu: Free lunch in diffusion u-net. *arXiv:2309.11497*, 2023. URL <https://arxiv.org/abs/2309.11497>.
- [39] Jiaming Song, Chenlin Meng, and Stefano Ermon. Denoising diffusion implicit models. *arXiv:2010.02502*, 2020. URL <https://arxiv.org/abs/2010.02502>.
- [40] Zhenxiong Tan, Songhua Liu, Xingyi Yang, Qiaochu Xue, and Xinchao Wang. Ominicontrol: Minimal and universal control for diffusion transformer. *arXiv:2411.15098*, 2024. URL <https://arxiv.org/abs/2411.15098>.
- [41] Gabriel Tjio, Jie Zhang, Xulei Yang, Yun Xing, Nhat Chung, Xiaofeng Cao, Ivor W. Tsang, Chee Keong Kwoh, and Qing Guo. Focus: Frequency-optimized conditioning of diffusion models for mitigating catastrophic forgetting during test-time adaptation. *arXiv:2508.14437*, 2025. URL <https://arxiv.org/abs/2508.14437>.
- [42] Andrey Voynov, Kfir Aberman, and Daniel Cohen-Or. Sketch-guided text-to-image diffusion models. In *SIGGRAPH 2023 Conference Proceedings*, pages 1–11, 2023. doi: 10.1145/3588432.3591560.
- [43] Xiaogang Wang and Xiaoou Tang. Face photo-sketch synthesis and recognition. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 31(11): 1955–1967, 2009. doi: 10.1109/TPAMI.2008.222.
- [44] Holger Winnemöller, Jan Eric Kyprianidis, and Sven C. Olsen. Xdog: An extended difference-of-gaussians compendium including advanced image stylization. *Computers & Graphics*, 36(6):740–753, 2012. doi: 10.1016/j.cag.2012.03.004.
- [45] Yifeng Xu, Zhenliang He, Shiguang Shan, and Xilin Chen. Ctrlora: An extensible and efficient framework for controllable image generation. *arXiv:2410.09400*, 2024. URL <https://arxiv.org/abs/2410.09400>.
- [46] Shuai Yang, Zhangyang Wang, Jiaying Liu, and Zongming Guo. Deep plastic surgery: Robust and controllable image editing with human-drawn sketches. In *Computer Vision – ECCV 2020*, pages 601–617, 2020. doi: 10.1007/978-3-030-58555-6_36.
- [47] Hu Ye, Jun Zhang, Sibio Liu, Xiao Han, and Wei Yang. Ip-adapter: Text compatible image prompt adapter for text-to-image diffusion models. *arXiv:2308.06721*, 2023. URL <https://arxiv.org/abs/2308.06721>.
- [48] Bo Yin, Xiaobin Hu, Xingyu Zhou, Yu He, Peng-Tao Jiang, Yue Liao, Junwei Zhu, Jiangning Zhang, Ying Tai, and Shuicheng Yan. Fera: Frequency-energy constrained routing for effective diffusion adaptation fine-tuning. *arXiv:2511.17979*, 2025. URL <https://arxiv.org/abs/2511.17979>.
- [49] Denis Zavadski, Johann-Friedrich Feiden, and Carsten Rother. Controlnet-xs: Rethinking the control of text-to-image diffusion models as feedback-control systems. *arXiv:2312.06573*, 2023. URL <https://arxiv.org/abs/2312.06573>.
- [50] Lvmin Zhang, Anyi Rao, and Maneesh Agrawala. Adding conditional control to text-to-image diffusion

models. *arXiv:2302.05543*, 2023. URL <https://arxiv.org/abs/2302.05543>.

- [51] Richard Zhang, Phillip Isola, Alexei A. Efros, Eli Shechtman, and Oliver Wang. The unreasonable effectiveness of deep features as a perceptual metric. *arXiv:1801.03924*, 2018. URL <https://arxiv.org/abs/1801.03924>.
- [52] Yuxuan Zhang, Yirui Yuan, Yiren Song, Haofan Wang, and Jiaming Liu. Easycontrol: Adding efficient and flexible control for diffusion transformer. *arXiv:2503.07027*, 2025. URL <https://arxiv.org/abs/2503.07027>.
- [53] Shihao Zhao, Dongdong Chen, Yen-Chun Chen, Jianmin Bao, Shaozhe Hao, Lu Yuan, and Kwan-Yee K. Wong. Uni-controlnet: All-in-one control to text-to-image diffusion models. *arXiv:2305.16322*, 2023. URL <https://arxiv.org/abs/2305.16322>.
- [54] Chao Zhou, Tianyi Wei, Yiling Chen, Wenbo Zhou, and Nenghai Yu. Rethinking multi-condition dits: Eliminating redundant attention via position-alignment and keyword-scoping. *arXiv:2602.06850*, 2026. URL <https://arxiv.org/abs/2602.06850>.